

Job Description

We are looking for an expert in machine learning to help us extract value from our data. You will lead all the processes from data collection, cleaning, and preprocessing, to training models and deploying them to production.

The ideal candidate will be passionate about artificial intelligence and stay up-to-date with the latest developments in the field.

Responsibilities

- Understanding business objectives and developing models that help to achieve them, along with metrics to track their progress
- Managing available resources such as hardware, data, and personnel so that deadlines are met
- Analyzing the ML algorithms that could be used to solve a given problem and ranking them by their success probability
- Exploring and visualizing data to gain an understanding of it, then identifying differences in data distribution that could affect performance when deploying the model in the real world
- Verifying data quality, and/or ensuring it via data cleaning
- Supervising the data acquisition process if more data is needed
- Finding available datasets online that could be used for training
- Defining validation strategies
- Defining the preprocessing or feature engineering to be done on a given dataset
- Defining data augmentation pipelines
- Training models and tuning their hyperparameters
- Analyzing the errors of the model and designing strategies to overcome them
- Deploying models to production

Skills

- Proficiency with a deep learning framework such as TensorFlow or Keras
- Proficiency with Python and basic libraries for machine learning such as scikit-learn and pandas
- Expertise in visualizing and manipulating big datasets
- Proficiency with OpenCV
- Familiarity with Linux
- Ability to select hardware to run an ML model with the required latency